

# Object Detection using Tensorflow

---

SAMIN PAYROSANGARI

SUPERVISER: DIEGO ESTEVES

01/02/2018



# Agenda

---


Introduction

Task

Requirments

Results

# CNN



# CNNs for object detection

---

Find specific objects in image and draw a bounding box

Extensions of image classification models

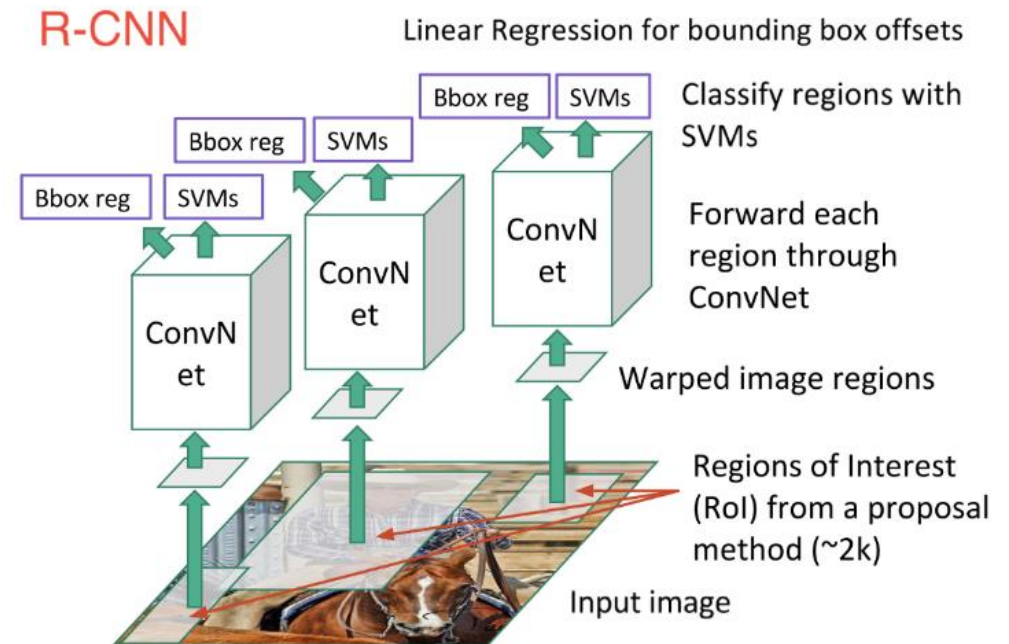
Models:

- **Faster R-CNN (R-CNN → Fast R-CNN → Faster R-CNN)**
- **R-FCN**
- **SSD**

# R-CNN

Region-based Convolutional Neural Network:

1. Selective search → Scan the input image for possible objects (~2000 **region proposals**)
2. Run **CNN** on top of each region
3. feed each output to:
  - a. an SVM to classify the region
  - b. a linear regressor to tighten the bounding box

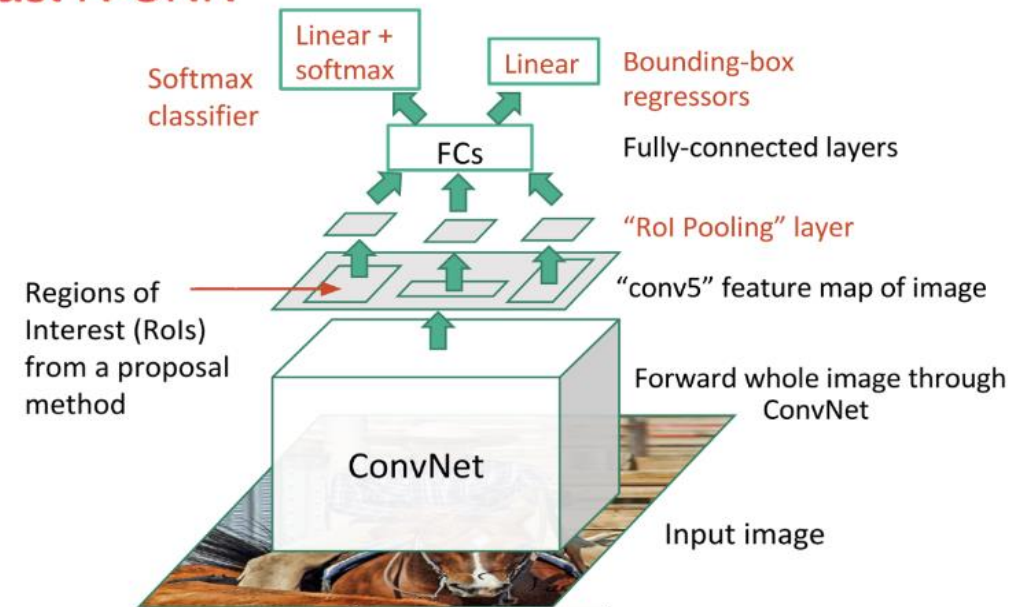


# Fast R-CNN

Improved on its detection speed:

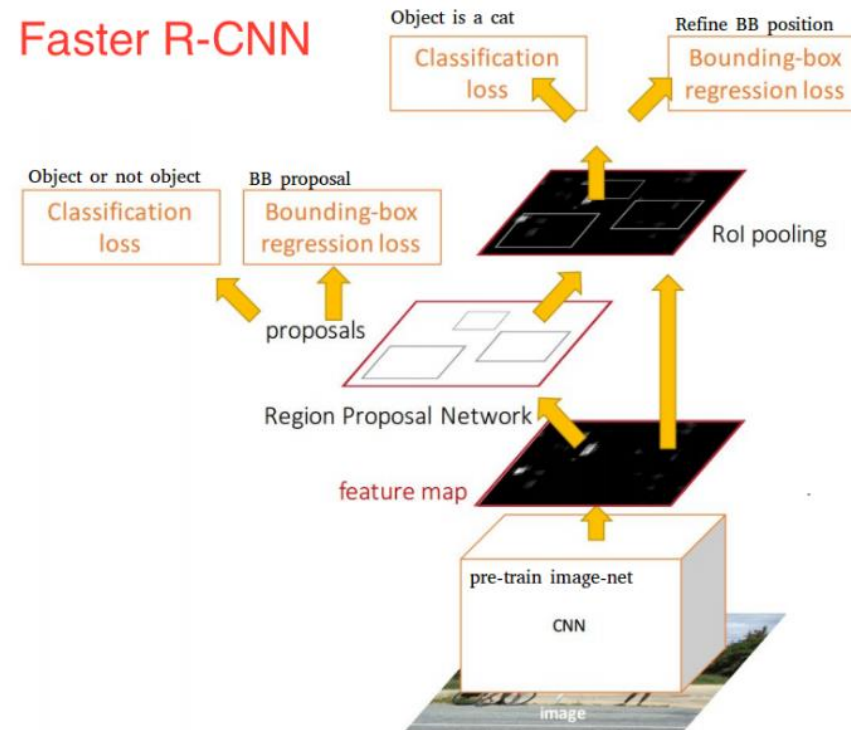
- Performing feature extraction over the image **before** proposing regions → running one CNN over the entire image instead
- Replacing the SVM with a softmax layer → extending the neural network for predictions instead of creating a new mode
- Region proposals based on the last feature map of CNN

## Fast R-CNN



# Faster R-CNN

Main insight → Replace the slow selective search algorithm with a fast neural net: **Region Proposal Network (RPN)** → **hypothesize object regions and then classify them.**



# RFCN (Region-based Fully Convolutional Net)

---

Motivation: Increase speed by maximizing shared computation

Fully convolutional

Shares 100% of the computations across every single output → using **position-sensitive score maps** → each score map activates if one specific part of an object is detected.



# RFCN simple explanation


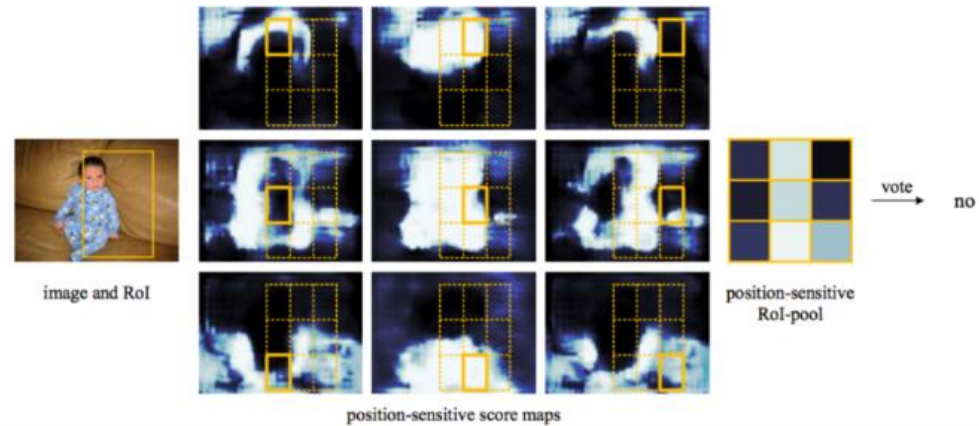


Figure 3: Visualization of R-FCN ( $k \times k = 3 \times 3$ ) for the *person* category.




# SSD (Single-Shot Detector)

---


Simultaneously predicting the bounding box and the class as it processes the image:

1. Goes through convolutional layers → yielding several sets of feature maps at different scales (e.g. 10x10, then 6x6, then 3x3, etc.)
2. For each location in *each* of these feature maps → considers small set of default bounding boxes.
3. For each box → **simultaneously** predict the bounding box offset and the class probabilities


# SSD simple example



(a) Image with GT boxes



(b)  $8 \times 8$  feature map



(c)  $4 \times 4$  feature map

# Task

---

Implement object detection models on specific datasets → adidas logo dataset, a dataset of 5 classes obtained from caltech 101 and scene 13 containing person, forest, mountain, highway and building.

Report the results and analysis

Compare the models performance

Solution: Google tensorflow api which comes with object detection models → retrain them with your own dataset.

# Requirments

---

Python 2.7 and 3.6

Tensorflow (python 3.6 compatible)

Object detection API library

Google cloud SDK (python 2.7 compatible)

Google could machine learning engine

# Work flow

---

Turn your dataset into TF records format:

1. Determine the coordinate of bounding boxex around desired objects in training set
2. Turn xml files of training and validation stes into 2 csv format file
3. Turn the csv file into one TF record
4. Generate a label Map for your dataset → Assign labels to int

# Workflow (continued)

---

Prepare google cloud platform project, storage and ML engine

Choose a model → Manipulate the config file of model:


- Config number of classes
- Config optimization and regularization parameters
- Config path to training TF record

Run the training process --> 100000 iterations


# Experiment 1: SSD on Adidas logo

High loss, low performance

Loss/classification\_loss/mul\_1



Loss/localization\_loss/mul\_1







# Experiment 2: SSD on 5 classes Data set

1. Adopted to train images for localization
2. Classification performance not so bad
3. Weak in localization


Loss/classification\_loss/mul\_1



Loss/localization\_loss/mul\_1



# Experiment 2: SSD on 5 classes Data set




# Experiment 3: RFCN on 5 classes dataset

---

1. Performs better than SSD
2. Better localization
3. Still weak
4. Wrong classification
5. Learned some classes better
6. Has not learned forest class

TotalLoss





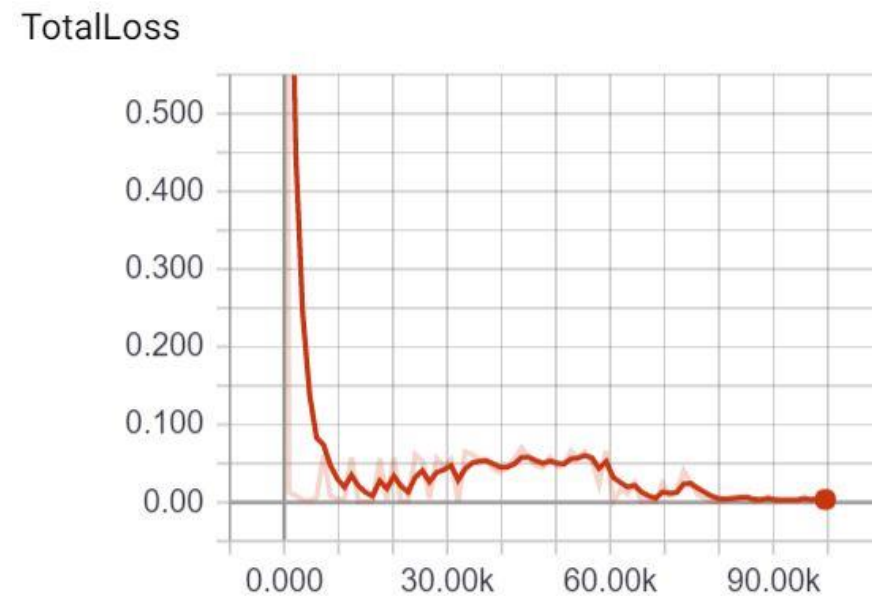
# Experiment 3: RFCN on 5 classes dataset



# Experiment 4: Faster R-CNN on 5 classes

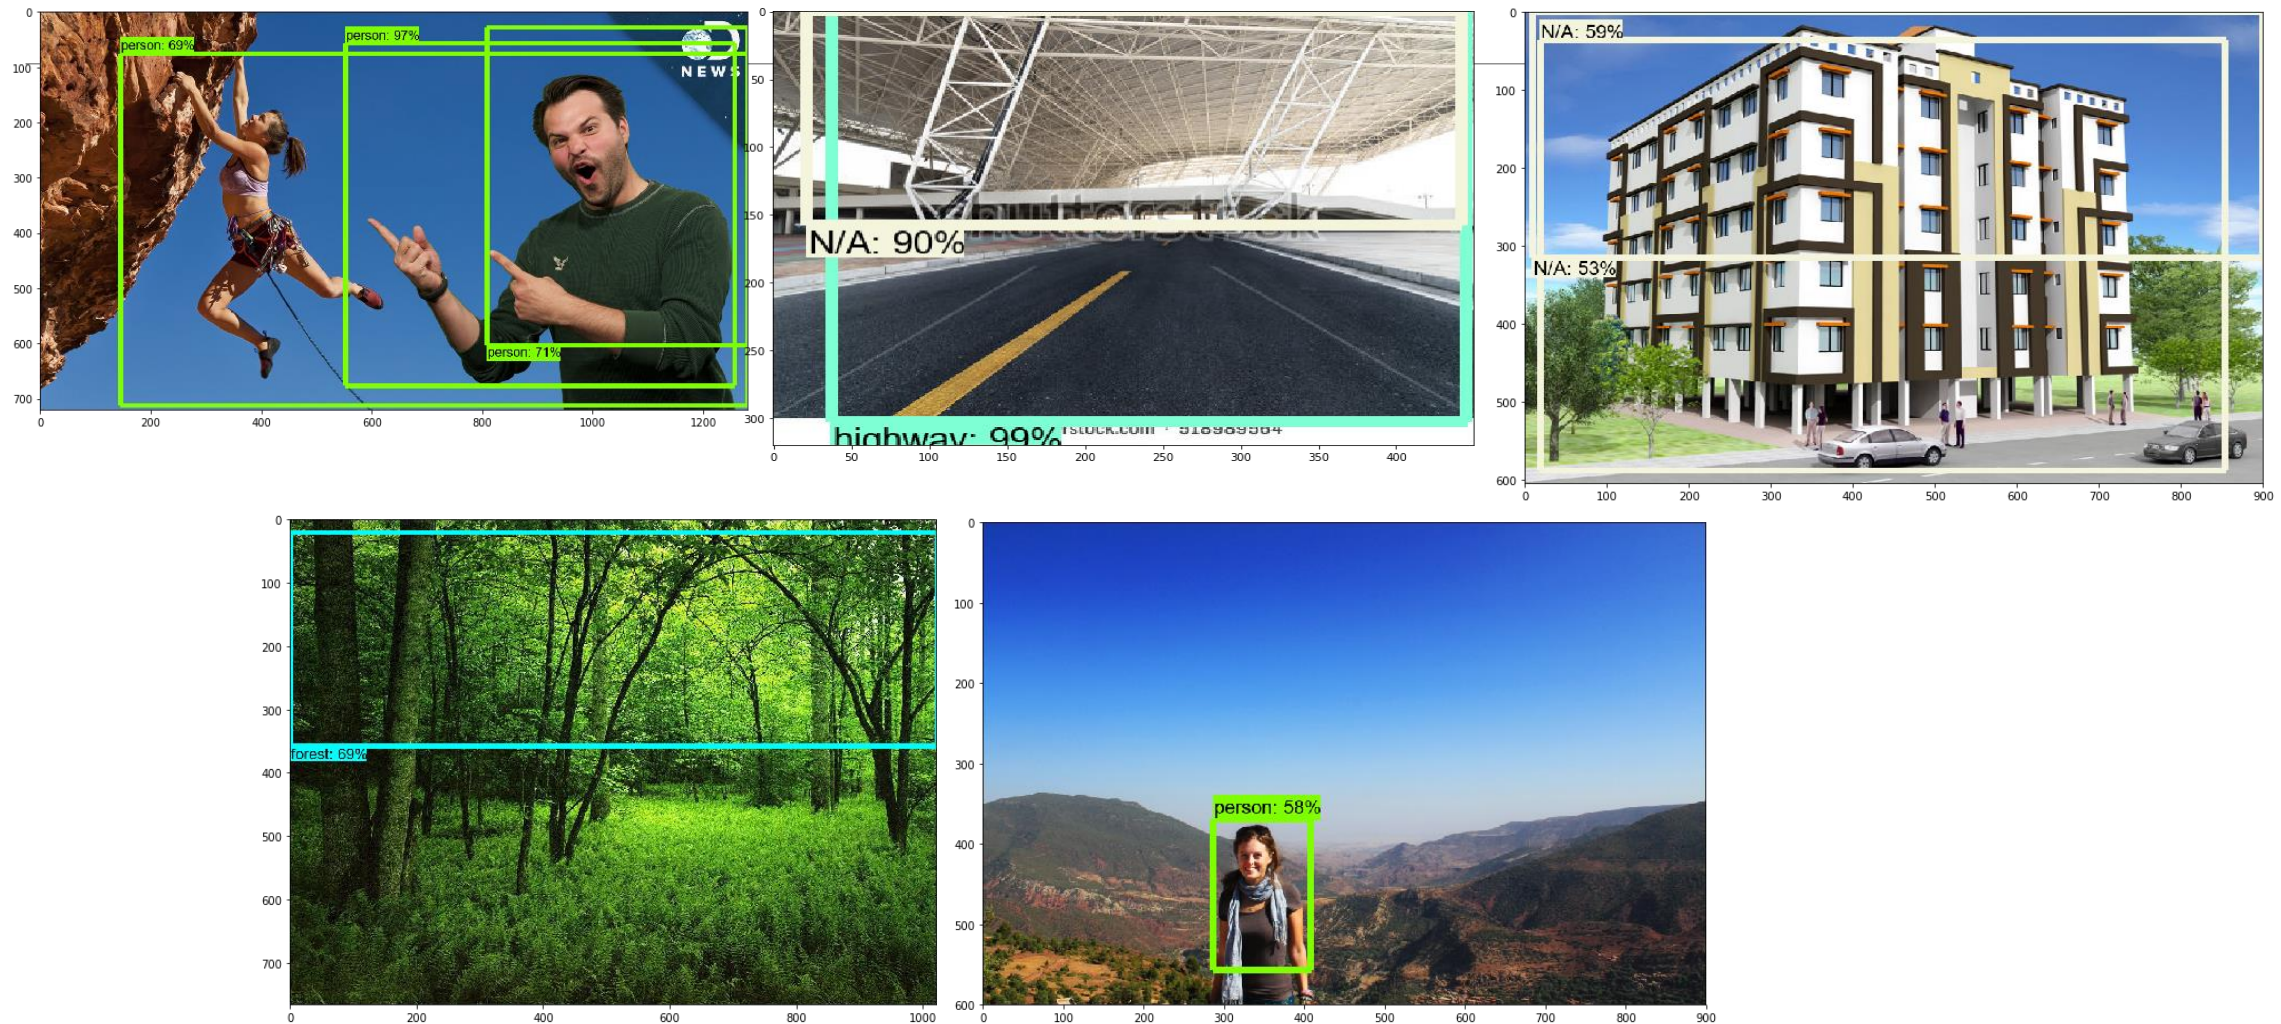
---

1. Good classification
2. Good localization
3. Still some wrong classification → Could improve by increasing the number of iterations





# Experiment 4: Faster R-CNN on 5 classes



# Conclusion

---

Ranking by performance:

1. Faster R-CNN
2. RFCN
3. SSD

Increasing the number of iterations for SSD does not seem to enhance the performance.

Increasing the number of iterations for RFCN might enhance the performance.

Faster R-CNN is slower but more powerful.

# References

---

- [1] <https://towardsdatascience.com/deep-learning-for-object-detection-a-comprehensive-review-73930816d8d9>
- [2] <https://medium.com/google-cloud/object-detection-tensorflow-and-google-cloud-platform-72e0a3f3bdd6>
- [3] <https://cloud.google.com/solutions/creating-object-detection-application-tensorflow>
- [4] <https://towardsdatascience.com/how-to-train-your-own-object-detector-with-tensorflows-object-detector-api-bec72ecfe1d9>
- [5] <https://github.com/tensorflow/models/issues/2739>